


Offload's advantage widens with length; the DRAM penalty does not



recompute/reload climbs 13.5→36.8x (recompute is superlinear, a block copy is linear), while reload/reuse decays 1.30→1.13x from its peak at 4096:

transfer is not the bottleneck. Grey axis is the control — all 7/7 points moved a full prefix, so no reload is a device hit in disguise.

reload/reuse is a lower bound and unevenly so (spurious decode block: 12.5% of reuse traffic at 1024, 0.2% at 65536) — trust the direction, not the slope.